

# ECGR\_4105\_HW2

October 9, 2022

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn import datasets
from sklearn.preprocessing import Normalizer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.metrics import classification_report
from sklearn.datasets import load_breast_cancer
```

```
[2]: diabset = pd.read_csv(r'C:\Users\homer\OneDrive\Documents\School_
↳Folder\diabetes.csv')
```

```
[3]: diabset.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: diabset.shape
```

```
[4]: (768, 9)
```

```

[5]: diab_index = diabset.index.values
    diab_index.shape

[5]: (768,)

[6]: diab_labels = np.reshape(diab_index,(768,1))

[7]: diab_data = np.concatenate([diabset,diab_labels],axis=1)

[8]: diab_data.shape

[8]: (768, 10)

[9]: diab_dataset = pd.DataFrame(diab_data)

[10]: diab_dataset.head()

[10]:
```

	0	1	2	3	4	5	6	7	8	9
0	6.0	148.0	72.0	35.0	0.0	33.6	0.627	50.0	1.0	0.0
1	1.0	85.0	66.0	29.0	0.0	26.6	0.351	31.0	0.0	1.0
2	8.0	183.0	64.0	0.0	0.0	23.3	0.672	32.0	1.0	2.0
3	1.0	89.0	66.0	23.0	94.0	28.1	0.167	21.0	0.0	3.0
4	0.0	137.0	40.0	35.0	168.0	43.1	2.288	33.0	1.0	4.0

```

[11]: diab_X = diab_dataset.values[:, 9]
    diab_Y = diab_dataset.values[:, 8]

[12]: diab_X_train, diab_X_test, diab_Y_train, diab_Y_test = train_test_split(diab_X,
    ↪diab_Y, test_size=0.2, random_state=42)

[13]: sc = StandardScaler()
    diab_X_reshape = diab_X_train.reshape(-1, 1)
    diab_X_std = sc.fit_transform(diab_X_reshape)
    diab_Xtest_reshape = diab_X_test.reshape(-1, 1)
    diab_Xtest_std = sc.transform(diab_Xtest_reshape)

[14]: diab_logreg = LogisticRegression(solver='liblinear', random_state=0)
    diab_logreg.fit(diab_X_std, diab_Y_train)

[14]: LogisticRegression(random_state=0, solver='liblinear')

[15]: diab_Y_pred = diab_logreg.predict(diab_Xtest_std)

[16]: diab_cnf_matrix = confusion_matrix(diab_Y_test, diab_Y_pred)
    diab_cnf_matrix

[16]: array([[99,  0],
    [55,  0]], dtype=int64)

```

```
[17]: print("Accuracy:", metrics.accuracy_score(diab_Y_test, diab_Y_pred))
      print("Precision:", metrics.precision_score(diab_Y_test, diab_Y_pred))
      print("Recall:", metrics.recall_score(diab_Y_test, diab_Y_pred))
```

Accuracy: 0.6428571428571429

Precision: 0.0

Recall: 0.0

C:\Users\homer\anaconda3\lib\site-


packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning:

Precision is ill-defined and being set to 0.0 due to no predicted samples. Use  
`zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

```
[18]: class_names = [0,1]
      fig, ax = plt.subplots()
      tick_marks = np.arange(len(class_names))
      plt.xticks(tick_marks, class_names)
      plt.yticks(tick_marks, class_names)
      sns.heatmap(pd.DataFrame(diab_cnf_matrix), annot=True, cmap="YlGnBu", fmt='g')
      ax.xaxis.set_label_position("top")
      plt.tight_layout()
      plt.title('Confusion matrix', y=1.1)
      plt.ylabel('Actual label')
      plt.xlabel('Predicted label')
```

```
[18]: Text(0.5, 257.44, 'Predicted label')
```



```
[19]: diab_kfold = KFold(n_splits=5, random_state=42, shuffle=True)
model = LogisticRegression(solver='liblinear')
results = cross_val_score(model, diab_X.reshape(-1, 1), diab_Y, cv=diab_kfold)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 65.106% (3.703%)

```
[20]: diab_kfold = KFold(n_splits=10, random_state=42, shuffle=True)
model = LogisticRegression(solver='liblinear')
results = cross_val_score(model, diab_X.reshape(-1, 1), diab_Y, cv=diab_kfold)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 65.096% (4.779%)

```
[21]: cancer = load_breast_cancer()
```

```
[22]: cancer_data = cancer.data
cancer_data.shape
```

```
[22]: (569, 30)
```

```
[23]: cancer_input = pd.DataFrame(cancer_data)
cancer_input.head()
```

```
[23]:
```

	0	1	2	3	4	5	6	7	8	\
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	

	9	...	20	21	22	23	24	25	26	27	\
0	0.07871	...	25.38	17.33	184.60	2019.0	0.1622	0.6656	0.7119	0.2654	
1	0.05667	...	24.99	23.41	158.80	1956.0	0.1238	0.1866	0.2416	0.1860	
2	0.05999	...	23.57	25.53	152.50	1709.0	0.1444	0.4245	0.4504	0.2430	
3	0.09744	...	14.91	26.50	98.87	567.7	0.2098	0.8663	0.6869	0.2575	
4	0.05883	...	22.54	16.67	152.20	1575.0	0.1374	0.2050	0.4000	0.1625	

	28	29
0	0.4601	0.11890
1	0.2750	0.08902
2	0.3613	0.08758
3	0.6638	0.17300
4	0.2364	0.07678

[5 rows x 30 columns]

```
[24]: cancer_labels = cancer.target
```

```
[25]: cancer_labels.shape
```

```
[25]: (569,)
```

```
[26]: labels = np.reshape(cancer_labels, (569,1))
```

```
[27]: final_cancer_data = np.concatenate([cancer_data,labels],axis=1)
```

```
[28]: final_cancer_data.shape
```

```
[28]: (569, 31)
```

```
[29]: cancer_dataset = pd.DataFrame(final_cancer_data)
```

```
[30]: features = cancer.feature_names
features
```

```
[30]: array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
'mean smoothness', 'mean compactness', 'mean concavity',
'mean concave points', 'mean symmetry', 'mean fractal dimension',
```

```
'radius error', 'texture error', 'perimeter error', 'area error',
'smoothness error', 'compactness error', 'concavity error',
'concave points error', 'symmetry error',
'fractal dimension error', 'worst radius', 'worst texture',
'worst perimeter', 'worst area', 'worst smoothness',
'worst compactness', 'worst concavity', 'worst concave points',
'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```

```
[31]: features_labels = np.append(features, 'label')
```

```
[32]: cancer_dataset.columns = features_labels
```

```
[33]: cancer_dataset.head()
```

```
[33]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
0	0.07871	...	17.33	184.60	2019.0	
1	0.05667	...	23.41	158.80	1956.0	
2	0.05999	...	25.53	152.50	1709.0	
3	0.09744	...	26.50	98.87	567.7	
4	0.05883	...	16.67	152.20	1575.0	

	worst smoothness	worst compactness	worst concavity	worst concave points	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

	worst symmetry	worst fractal dimension	label
0	0.4601	0.11890	0.0
1	0.2750	0.08902	0.0
2	0.3613	0.08758	0.0
3	0.6638	0.17300	0.0

4                    0.2364                    0.07678           0.0

[5 rows x 31 columns]

```
[34]: cancer_dataset['label'].replace(0, 'Benign', inplace=True)
      cancer_dataset['label'].replace(1, 'Malignant', inplace=True)
```

```
[35]: cancer_dataset.tail()
```

```
[35]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
564	21.56	22.39	142.00	1479.0	0.11100	
565	20.13	28.25	131.20	1261.0	0.09780	
566	16.60	28.08	108.30	858.1	0.08455	
567	20.60	29.33	140.10	1265.0	0.11780	
568	7.76	24.54	47.92	181.0	0.05263	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
564	0.11590	0.24390	0.13890	0.1726	
565	0.10340	0.14400	0.09791	0.1752	
566	0.10230	0.09251	0.05302	0.1590	
567	0.27700	0.35140	0.15200	0.2397	
568	0.04362	0.00000	0.00000	0.1587	

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
564	0.05623	...	26.40	166.10	2027.0	
565	0.05533	...	38.25	155.00	1731.0	
566	0.05648	...	34.12	126.70	1124.0	
567	0.07016	...	39.42	184.60	1821.0	
568	0.05884	...	30.37	59.16	268.6	

	worst smoothness	worst compactness	worst concavity	\
564	0.14100	0.21130	0.4107	
565	0.11660	0.19220	0.3215	
566	0.11390	0.30940	0.3403	
567	0.16500	0.86810	0.9387	
568	0.08996	0.06444	0.0000	

	worst concave points	worst symmetry	worst fractal dimension	label
564	0.2216	0.2060	0.07115	Benign
565	0.1628	0.2572	0.06637	Benign
566	0.1418	0.2218	0.07820	Benign
567	0.2650	0.4087	0.12400	Benign
568	0.0000	0.2871	0.07039	Malignant

[5 rows x 31 columns]

```
[36]: cancer_X = cancer_dataset.iloc[:,0:29].values
      cancer_Y = cancer_dataset.iloc[:,30].values
```

```
[37]: cancer_X_train, cancer_X_test, cancer_Y_train, cancer_Y_test =
      ↪train_test_split(cancer_X, cancer_Y, test_size=0.2, random_state=42)
```

```
[38]: sc_X = StandardScaler()
      cancer_X_trainstd = sc_X.fit_transform(cancer_X_train)
      cancer_X_teststd = sc_X.transform(cancer_X_test)
```

```
[39]: cancerClass = LogisticRegression(random_state=42)
      cancerClass.fit(cancer_X_trainstd, cancer_Y_train)
```

```
[39]: LogisticRegression(random_state=42)
```

```
[40]: cancer_Y_pred = cancerClass.predict(cancer_X_teststd)
```

```
[41]: cancer_Y_pred[0:9]
```

```
[41]: array(['Malignant', 'Benign', 'Benign', 'Malignant', 'Malignant',
      'Benign', 'Benign', 'Benign', 'Malignant'], dtype=object)
```

```
[42]: cancer_cnf_matrix = confusion_matrix(cancer_Y_test, cancer_Y_pred)
      cancer_cnf_matrix
```

```
[42]: array([[42,  1],
      [ 1, 70]], dtype=int64)
```

```
[43]: print("Accuracy:",metrics.accuracy_score(cancer_Y_test, cancer_Y_pred))
      print("Precision:",metrics.precision_score(cancer_Y_test, cancer_Y_pred,
      ↪pos_label="Benign"))
      print("Recall:",metrics.recall_score(cancer_Y_test, cancer_Y_pred,
      ↪pos_label="Benign"))
```

Accuracy: 0.9824561403508771

Precision: 0.9767441860465116


Recall: 0.9767441860465116

```
[44]: class_names = [0,1]
      fig, ax = plt.subplots()
      tick_marks = np.arange(len(class_names))
      plt.xticks(tick_marks, class_names)
      plt.yticks(tick_marks, class_names)
      sns.heatmap(pd.DataFrame(cancer_cnf_matrix), annot=True, cmap="YlGnBu", fmt='g')
      ax.xaxis.set_label_position("top")
      plt.tight_layout()
      plt.title('Confusion matrix', y=1.1)
      plt.ylabel('Actual label')
```



```
plt.xlabel('Predicted label')
```

```
[44]: Text(0.5, 257.44, 'Predicted label')
```



```
[48]: C = [10, 1, .1, .001]
for c in C:
    clf = LogisticRegression(penalty='l1', C=c, solver='liblinear')
    clf.fit(cancer_X_trainstd, cancer_Y_train)
    print('C:', c)
    print('Training accuracy: ', clf.score(cancer_X_trainstd, cancer_Y_train))
    print('Test accuracy: ', clf.score(cancer_X_teststd, cancer_Y_test))
    print(' ')
```

```
C: 10
Training accuracy:  0.9912087912087912
Test accuracy:    0.956140350877193
```

```
C: 1
Training accuracy:  0.989010989010989
Test accuracy:     0.9736842105263158
```

C: 0.1  
Training accuracy: 0.9802197802197802  
Test accuracy: 0.9649122807017544

C: 0.001  
Training accuracy: 0.37142857142857144  
Test accuracy: 0.37719298245614036

```
[49]: cancer_kfold = KFold(n_splits=5, random_state=42, shuffle=True)
model = LogisticRegression(solver='liblinear')
results = cross_val_score(model, cancer_X, cancer_Y, cv=cancer_kfold)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 94.723% (2.675%)

```
[50]: cancer_kfold = KFold(n_splits=10, random_state=42, shuffle=True)
model = LogisticRegression(solver='liblinear')
results = cross_val_score(model, cancer_X, cancer_Y, cv=cancer_kfold)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 94.897% (3.477%)

```
[51]: for c in C:
    cancer_kfold = KFold(n_splits=5, random_state=42, shuffle=True)
    model = LogisticRegression(penalty='l1', C=c, solver='liblinear')
    results = cross_val_score(model, cancer_X, cancer_Y, cv=cancer_kfold)
    print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.
    ↪std()*100.0))
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:  
ConvergenceWarning: Liblinear failed to converge, increase the number of  
iterations.  
warnings.warn(

Accuracy: 96.479% (2.097%)

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:  
ConvergenceWarning: Liblinear failed to converge, increase the number of  
iterations.  
warnings.warn(

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:  
ConvergenceWarning: Liblinear failed to converge, increase the number of  
iterations.  
warnings.warn(

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:  
ConvergenceWarning: Liblinear failed to converge, increase the number of  
iterations.  
warnings.warn(

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

Accuracy: 94.898% (3.219%)

Accuracy: 93.316% (3.760%)

Accuracy: 91.559% (3.039%)

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

```
[52]: for c in C:
        cancer_kfold = KFold(n_splits=10, random_state=42, shuffle=True)
        model = LogisticRegression(penalty='l1', C=c, solver='liblinear')
        results = cross_val_score(model, cancer_X, cancer_Y, cv=cancer_kfold)
        print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.
        ↪std()*100.0))
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

Accuracy: 95.949% (2.966%)

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```
warnings.warn(
```

C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\\_base.py:1206:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

```

warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
Accuracy: 94.721% (3.346%)
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(
C:\Users\homer\anaconda3\lib\site-packages\sklearn\svm\_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.

```

```
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
```

```
warnings.warn(
```

```
Accuracy: 93.318% (4.636%)
```

```
Accuracy: 91.557% (5.090%)
```

```
[ ]:
```